

# Clear and Compress: Computing Persistent Homology in Chunks


**Ulrich Bauer**<sup>1</sup>   Michael Kerber<sup>2</sup>   Jan Reininghaus<sup>1</sup>

<sup>1</sup> Institute of Science and Technology (IST) Austria

<sup>2</sup>Stanford University and Max Planck Center for Visual Computing and  
Communication, Saarbrücken, Germany

SIAM AG 2013

# Overview



## Complexity analysis


$$O(m\ell^3 + g\ell n + g^3)$$

- ▶  $n$  = #simplices
- ▶  $\ell$  = max. size of chunk
- ▶  $m$  = #chunks
- ▶  $g$  = #non-local pairs

## Experimental results

- ▶ Parallelized version
- ▶ Outperforms standard algorithm, Dionysus
- ▶ Faster than clear optimization

# Standard Persistence Algorithm (reminder)




	1	2	3	4	5	6	7
1			1		1		
2			1			1	
3							1
4					1	1	
5							1
6							1
7							

## Algorithm:

for  $i$  from 1 to  $n$ :

- ▶ while  $\text{pivot}[j] = \text{pivot}[i]$  for some  $j < i$ :  
add column  $j$  to column  $i$

# Standard Persistence Algorithm (reminder)




	1	2	3	4	5	6	7
1			1		1		
2			1			1	
3							1
4					1	1	
5							1
6							1
7							

## Algorithm:

for  $i$  from 1 to  $n$ :

- ▶ while  $\text{pivot}[j] = \text{pivot}[i]$  for some  $j < i$ :  
add column  $j$  to column  $i$

# Standard Persistence Algorithm (reminder)




	1	2	3	4	5	6	7
1			1		1	1	
2			1			1	
3							1
4					1	0	
5							1
6							1
7							

## Algorithm:

for  $i$  from 1 to  $n$ :

- ▶ while  $\text{pivot}[j] = \text{pivot}[i]$  for some  $j < i$ :  
add column  $j$  to column  $i$

# Standard Persistence Algorithm (reminder)




	1	2	3	4	5	6	7
1			1		1	0	
2			1			0	
3							1
4					1		
5							1
6							1
7							

## Algorithm:

for  $i$  from 1 to  $n$ :

- ▶ while  $\text{pivot}[j] = \text{pivot}[i]$  for some  $j < i$ :  
add column  $j$  to column  $i$

# Standard Persistence Algorithm (reminder)




	1	2	3	4	5	6	7
1			1		1		
2			1				
3							1
4					1		
5							1
6							1
7							

## Algorithm:

for  $i$  from 1 to  $n$ :

- ▶ while  $\text{pivot}[j] = \text{pivot}[i]$  for some  $j < i$ :  
add column  $j$  to column  $i$

# Standard Persistence Algorithm (reminder)



	1	2	3	4	5	6	7
1			1		1		
2			1				
3							1
4					1		
5							1
6							1
7							

- ▶ positive inessential, negative, and essential simplices/columns
- ▶ Column  $i$  is zero  $\Leftrightarrow i$ -th simplex is positive
- ▶ All column additions within the same dimension
- ▶ Pivots do not depend on the order of operations



# Clear Optimization [Chen, K. 2011]

	1	2	3	4	5	6	7
1			1			1	
2			1		1		
3							1
4					1	1	
5							1
6							1
7							

$(i, j)$  is persistence pair

$\Rightarrow i$ -th simplex is positive

$\Rightarrow$  column  $i$  reduces to 0

# Clear Optimization [Chen, K. 2011]

	1	2	3	4	5	6	7
1			1			1	
2			1		1		
3							1
4					1	1	
5							1
6							1
7							

## Algorithm:

for  $\delta$  from  $\dim K$  downto 0:

- ▶ For simplices of dim.  $\delta$ , left to right:
  - ▶ while  $\text{pivot}[j]=\text{pivot}[i]$  for some  $j < i$ , add column  $j$  to column  $i$
- ▶ For any column  $i$  with  $\text{pivot}[i]=j$ :  
set column  $j$  to zero

# Compress Optimization [Zomorodian, Carlsson 2004]

	1	2	3	4	5	6	7	8	9	10	11	12	13
1			1			1			1				
2			1		1								
3							1			1		1	1
4					1	1							
5							1			1			1
6							1					1	1
7													
8									1				
9										1		1	

$(i, j)$  is persistence pair

$\Rightarrow j$  is not the pivot of any column

$\Rightarrow$  Setting row  $j$  to zero is fine!

# Compress Optimization [Zomorodian, Carlsson 2004]

	1	2	3	4	5	6	7	8	9	10	11	12	13
1			1			1			1				
2			1		1								
3							1			1		1	1
4					1	1							
5							1			1			1
6							1					1	1
7													
8									1				
9										1		1	

## Algorithm:


for  $i$  from 1 to  $n$ :

- ▶ Remove row indices for negative simplices (compression)
- ▶ while  $\text{pivot}[j] = \text{pivot}[i]$  for some  $j < i$ ,  
add column  $j$  to column  $i$

# Chunk Algorithm: Local Reduction

*Local simplex*: Paired with simplex in same or adjacent chunk

*Global simplex*: non-local (must have persistence  $\geq \ell$ )



Local reduction algorithm:

- ▶ Reduce columns using only columns from same chunk
- ▶ Reduce columns using only columns from same or left-neighboring chunk

Main observation:

- ▶ Local columns are reduced

Running time:  $O(m\ell^3)$

## Chunk Algorithm: Local Reduction

*Local simplex*: Paired with simplex in same or adjacent chunk

*Global simplex*: non-local (must have persistence  $\geq \ell$ )

[illegible]

### Local reduction algorithm:


- ▶ Reduce columns using only columns from same chunk
- ▶ Reduce columns using only columns from same or left-neighboring chunk

Main observation:

- ▶ Local columns are reduced

Running time:  $O(m\ell^3)$

# Chunk Algorithm: Compression




For each global column:  
iterate over entries from bottom to top

- ▶ Global index: skip
- ▶ Local positive: add (local) column with same pivot
- ▶ Local negative: set to zero (compress)

Running time:  $O(g(n + n\ell)) = O(gn\ell)$

Result: Only global indices remain

# Chunk Algorithm: Compression



For each global column:  
iterate over entries from bottom to top


- ▶ Global index: skip
- ▶ Local positive: add (local) column with same pivot
- ▶ Local negative: set to zero (compress)

Running time:  $O(g(n + n\ell)) = O(gn\ell)$

Result: Only global indices remain



# Chunk Algorithm: Compression



For each global column:


iterate over entries from bottom to top

- ▶ Global index: skip
- ▶ Local positive: add (local) column with same pivot
- ▶ Local negative: set to zero (compress)

Running time:  $O(g(n + n\ell)) = O(gn\ell)$

Result: Only global indices remain

# Chunk Algorithm: Compression




For each global column:  
iterate over entries from bottom to top

- ▶ Global index: skip
- ▶ Local positive: add (local) column with same pivot
- ▶ Local negative: set to zero (compress)

Running time:  $O(g(n + n\ell)) = O(gn\ell)$

Result: Only global indices remain

# Chunk Algorithm: Compression




For each global column:  
iterate over entries from bottom to top

- ▶ Global index: skip
- ▶ Local positive: add (local) column with same pivot
- ▶ Local negative: set to zero (compress)

Running time:  $O(g(n + n\ell)) = O(gn\ell)$

Result: Only global indices remain

# Chunk Algorithm: Compression



For each global column:


iterate over entries from bottom to top

- ▶ Global index: skip
- ▶ Local positive: add (local) column with same pivot
- ▶ Local negative: set to zero (compress)

Running time:  $O(g(n + n\ell)) = O(gn\ell)$

Result: Only global indices remain

# Chunk Algorithm: Compression



For each global column:


iterate over entries from bottom to top

- ▶ Global index: skip
- ▶ Local positive: add (local) column with same pivot
- ▶ Local negative: set to zero (compress)

Running time:  $O(g(n + n\ell)) = O(gn\ell)$

Result: Only global indices remain

# Chunk Algorithm: Global Reduction



Reduce  $g \times g$  submatrix:  $O(g^3)$

Further optimizations:

- ▶ Avoid additions during compression
  - ▶ *inactive indices*: either
    - ▶ negative, or
    - ▶ positive and pivot of column with otherwise only inactive entries
  - ▶ inactive indices can be set to zero
- ▶ Avoid compression of global positive columns (clear optimization)

# Complexity Analysis

$$O(m\ell^3 + gn\ell + g^3)$$

- ▶  $n = \# \text{simplices}$
- ▶  $\ell = \text{max. size of chunk}$
- ▶  $m = \# \text{chunks}$
- ▶  $g = \# \text{non-local pairs}$

$\sqrt{n}$  chunks of size  $\sqrt{n}$ :

$$O(n^2 + g_1 n \sqrt{n} + g_1^3)$$


$\frac{n}{\log n}$  chunks of size  $\log n$ :

$$O(n \log^2 n + g_2 n \log n + g_2^3)$$


Special case:  $d$ -dimensional image + lower star filtration:

$$O(n + gn + g^3) = O(gn + g^3)$$


# Experimental results



*Smooth*




*Noise*



*Mixed*

Morse filtrations



*Torus*

Alpha filtration



*Mumford*

Rips filtration




# Experimental results

Dataset	$n \cdot 10^{-6}$	std. [8]	twist [4]	cohom. [5]	DMT [10]	$g/n$	chunk (1x)	chunk (12x)
Smooth	16.6	383s	3.1s	65.8s	2.0s	0%	5.0s	0.9s
Smooth $^\perp$	16.6	432s	11.3s	20.8s	–	0%	6.3s	0.9s
Noise	16.6	336s	17.2s	15971s	13.0s	9%	28.3s	6.3s
Noise $^\perp$	16.6	1200s	29.0s	190.1s	–	9%	31.1s	5.8s
Mixed	16.6	330s	5.8s	50927s	12.3s	5%	21.6s	2.4s
Mixed $^\perp$	16.6	446s	13.0s	32.7s	–	5%	32.0s	2.9s
Torus	0.6	52s	0.3s	1.6s	–	7%	0.3s	0.1s
Torus $^\perp$	0.6	24s	0.3s	1.4s	–	7%	0.9s	0.2s
Mumford	2.4	38s	35.2s	2.8s	–	82%	14.6s	1.8s
Mumford $^\perp$	2.4	58s	0.2s	184.1s	–	82%	1.5s	0.4s

<http://phat.googlecode.com>

# Summary



- ▶ Homological version of discrete Morse theory approach  
[Günther, R., Wagner, Hotz 2012]
- ▶ Parallelizable
- ▶ Distributed memory?
- ▶ Streaming?