

Ripser

Efficient Computation of Vietoris–Rips Persistence Barcodes

Ulrich Bauer

TUM

July 18, 2017

Workshop Topology, Computation, and Data Analysis
Schloss Dagstuhl

Vietoris–Rips persistence

Vietoris–Rips filtrations

Consider a finite metric space (X, d) .

The *Vietoris–Rips complex* is the simplicial complex

$$\text{Rips}_t(X) = \{S \subseteq X \mid \text{diam } S \leq t\}$$

- 1-skeleton: all edges with pairwise distance $\leq t$
- all possible higher simplices (flag complex)

Vietoris–Rips filtrations

Consider a finite metric space (X, d) .

The *Vietoris–Rips complex* is the simplicial complex

$$\text{Rips}_t(X) = \{S \subseteq X \mid \text{diam } S \leq t\}$$

- 1-skeleton: all edges with pairwise distance $\leq t$
- all possible higher simplices (flag complex)

Goal:

- compute persistence barcodes for $H_d(\text{Rips}_t(X))$
(in dimensions $0 \leq d \leq k$)

Demo: Ripser

Example data set:

- 192 points on \mathbb{S}^2
- persistent homology barcodes up to dimension 2
- over 56 mio. simplices in 3-skeleton

Demo: Ripser

Example data set:

- 192 points on \mathbb{S}^2
- persistent homology barcodes up to dimension 2
- over 56 mio. simplices in 3-skeleton

Comparison with other software:

- javaplex: 3200 seconds, 12 GB
- Dionysus: 533 seconds, 3.4 GB
- GUDHI: 75 seconds, 2.9 GB
- DIPHA: 50 seconds, 6 GB
- Eirene: 12 seconds, 1.5 GB

Demo: Ripser

Example data set:

- 192 points on \mathbb{S}^2
- persistent homology barcodes up to dimension 2
- over 56 mio. simplices in 3-skeleton

Comparison with other software:

- javaplex: 3200 seconds, 12 GB
- Dionysus: 533 seconds, 3.4 GB
- GUDHI: 75 seconds, 2.9 GB
- DIPHA: 50 seconds, 6 GB
- Eirene: 12 seconds, 1.5 GB

Ripser: 1.2 seconds, 152 MB

Ripser

A software for computing Vietoris–Rips persistence barcodes

- about 1000 lines of C++ code, no external dependencies
- support for
 - coefficients in a prime field \mathbb{F}_p
 - sparse distance matrices for distance threshold
- open source (<http://ripser.org>)
 - released in July 2016
- online version (<http://live.ripser.org>)
 - launched in August 2016
- most efficient software for Vietoris–Rips persistence
 - computes H^2 barcode for 50 000 random points on a torus in 136 seconds / 9 GB (using distance threshold)
- 2016 ATMCS Best New Software Award (jointly with RIVET)

Design goals

Goals for previous projects:

- PHAT [B, Kerber, Reininghaus, Wagner 2013]:
fast persistence computation (matrix reduction only)
- DIPHA [B, Kerber, Reininghaus 2014]:
distributed persistence computation

Design goals

Goals for previous projects:

- PHAT [B, Kerber, Reininghaus, Wagner 2013]:
fast persistence computation (matrix reduction only)
- DIPHA [B, Kerber, Reininghaus 2014]:
distributed persistence computation

Goals for Ripser:

- Use as little memory as possible
- Be reasonable about computation time

The four special ingredients

The improved performance is based on 4 insights:

- Clearing inessential columns [Chen, Kerber 2011]
- Computing cohomology [de Silva et al. 2011]
- Implicit matrix reduction
- Apparent and emergent pairs

The four special ingredients

The improved performance is based on 4 insights:

- Clearing inessential columns [Chen, Kerber 2011]
- Computing cohomology [de Silva et al. 2011]
- Implicit matrix reduction
- Apparent and emergent pairs

Lessons from PHAT:

- Clearing and cohomology yield considerable speedup,
- but only when *both* are used in conjunction!

Matrix reduction

Matrix reduction algorithm

Setting:

- finite metric space X , n points
- persistent homology $H_d(\text{Rips}_t(X); \mathbb{F}_2)$ in dimensions $d \leq k$

Notation:

- D : boundary matrix of filtration
- R_i : i th column of R

Matrix reduction algorithm

Setting:

- finite metric space X , n points
- persistent homology $H_d(\text{Rips}_t(X); \mathbb{F}_2)$ in dimensions $d \leq k$

Notation:

- D : boundary matrix of filtration
- R_i : i th column of R

Algorithm:

- $R = D, V = I$
- while $\exists i < j$ with $\text{pivot } R_i = \text{pivot } R_j$
 - add R_i to R_j , add V_i to V_j

Matrix reduction algorithm

Setting:

- finite metric space X , n points
- persistent homology $H_d(\text{Rips}_t(X); \mathbb{F}_2)$ in dimensions $d \leq k$

Notation:

- D : boundary matrix of filtration
- R_i : i th column of R

Algorithm:

- $R = D, V = I$
- while $\exists i < j$ with $\text{pivot } R_i = \text{pivot } R_j$
 - add R_i to R_j , add V_i to V_j

Result:

- $R = D \cdot V$ is reduced (unique pivots)
- V is full rank upper triangular

Compatible basis cycles

For a reduced boundary matrix $R = D \cdot V$, call

$P = \{i : R_i = 0\}$	<i>positive</i> indices,
$N = \{j : R_j \neq 0\}$	<i>negative</i> indices,
$E = P \setminus \text{pivots } R$	<i>essential</i> indices.

Then

Compatible basis cycles

For a reduced boundary matrix $R = D \cdot V$, call

$$\begin{array}{ll} P = \{i : R_i = 0\} & \text{positive indices,} \\ N = \{j : R_j \neq 0\} & \text{negative indices,} \\ E = P \setminus \text{pivots } R & \text{essential indices.} \end{array}$$

Then

$$\widetilde{\Sigma}_Z = \{V_i \mid i \in P\} \quad \text{is a basis of } Z_*,$$

Compatible basis cycles

For a reduced boundary matrix $R = D \cdot V$, call

$P = \{i : R_i = 0\}$	<i>positive</i> indices,
$N = \{j : R_j \neq 0\}$	<i>negative</i> indices,
$E = P \setminus \text{pivots } R$	<i>essential</i> indices.

Then

$\widetilde{\Sigma}_Z = \{V_i \mid i \in P\}$	is a basis of Z_* ,
$\Sigma_B = \{R_j \mid j \in N\}$	is a basis of B_* ,

Compatible basis cycles

For a reduced boundary matrix $R = D \cdot V$, call

$P = \{i : R_i = 0\}$	<i>positive</i> indices,
$N = \{j : R_j \neq 0\}$	<i>negative</i> indices,
$E = P \setminus \text{pivots } R$	<i>essential</i> indices.

Then

$\widetilde{\Sigma}_Z = \{V_i \mid i \in P\}$	is a basis of Z_* ,
$\Sigma_B = \{R_j \mid j \in N\}$	is a basis of B_* ,
$\Sigma_Z = \Sigma_B \cup \{V_i \mid i \in E\}$	is <i>another</i> basis of Z_* .

Compatible basis cycles

For a reduced boundary matrix $R = D \cdot V$, call

$$\begin{array}{ll} P = \{i : R_i = 0\} & \text{positive indices,} \\ N = \{j : R_j \neq 0\} & \text{negative indices,} \\ E = P \setminus \text{pivots } R & \text{essential indices.} \end{array}$$

Then

$$\begin{array}{ll} \widetilde{\Sigma}_Z = \{V_i \mid i \in P\} & \text{is a basis of } Z_*, \\ \Sigma_B = \{R_j \mid j \in N\} & \text{is a basis of } B_*, \\ \Sigma_Z = \Sigma_B \cup \{V_i \mid i \in E\} & \text{is another basis of } Z_*. \end{array}$$

Persistent homology is generated by the basis cycles Σ_Z .

Compatible basis cycles

For a reduced boundary matrix $R = D \cdot V$, call

$$\begin{aligned} P &= \{i : R_i = 0\} && \text{positive indices,} \\ N &= \{j : R_j \neq 0\} && \text{negative indices,} \\ E &= P \setminus \text{pivots } R && \text{essential indices.} \end{aligned}$$

Then

$$\begin{aligned} \widetilde{\Sigma}_Z &= \{V_i \mid i \in P\} && \text{is a basis of } Z_*, \\ \Sigma_B &= \{R_j \mid j \in N\} && \text{is a basis of } B_*, \\ \Sigma_Z &= \Sigma_B \cup \{V_i \mid i \in E\} && \text{is another basis of } Z_*. \end{aligned}$$

Persistent homology is generated by the basis cycles Σ_Z .

- Persistence intervals: $\{[i, j) \mid i = \text{pivot } R_j\} \cup \{[i, \infty) \mid i \in E\}$
- Columns with non-essential positive indices never used!

Clearing

Clearing non-essential positive columns

Idea [Chen, Kerber 2011]:

- Don't reduce at non-essential positive indices
- Reduce boundary matrices of $\partial_d : C_d \rightarrow C_{d-1}$ in decreasing dimension $d = k + 1, \dots, 1$
- Whenever $i = \text{pivot } R_j$ (in matrix for ∂_d)
 - Set R_i to 0 (in matrix for ∂_{d-1})

Clearing non-essential positive columns

Idea [Chen, Kerber 2011]:

- Don't reduce at non-essential positive indices
- Reduce boundary matrices of $\partial_d : C_d \rightarrow C_{d-1}$ in decreasing dimension $d = k + 1, \dots, 1$
- Whenever $i = \text{pivot } R_j$ (in matrix for ∂_d)
 - Set R_i to 0 (in matrix for ∂_{d-1})
 - Set V_i to R_j

Clearing non-essential positive columns

Idea [Chen, Kerber 2011]:

- Don't reduce at non-essential positive indices
- Reduce boundary matrices of $\partial_d : C_d \rightarrow C_{d-1}$ in decreasing dimension $d = k + 1, \dots, 1$
- Whenever $i = \text{pivot } R_j$ (in matrix for ∂_d)
 - Set R_i to 0 (in matrix for ∂_{d-1})
 - Set V_i to R_j
- Still yields $R = D \cdot V$ reduced, V full rank upper triangular

Clearing non-essential positive columns

Idea [Chen, Kerber 2011]:

- Don't reduce at non-essential positive indices
- Reduce boundary matrices of $\partial_d : C_d \rightarrow C_{d-1}$ in decreasing dimension $d = k + 1, \dots, 1$
- Whenever $i = \text{pivot } R_j$ (in matrix for ∂_d)
 - Set R_i to 0 (in matrix for ∂_{d-1})
 - Set V_i to R_j
- Still yields $R = D \cdot V$ reduced, V full rank upper triangular

Note:

- reducing *positive* columns typically harder than negative
- with clearing: need only reduce *essential* positive columns

Cohomology

Persistent cohomology

We have seen: many columns of $R = D \cdot V$ are not needed

- Skip those inessential columns in matrix reduction

Persistent cohomology

We have seen: many columns of $R = D \cdot V$ are not needed

- Skip those inessential columns in matrix reduction

For persistence barcodes in low dimensions $d \leq k$:

- Number of skipped indices for reducing D^T (cohomology) is much larger than for D (homology)
 - reducing boundary matrix produces basis for $H_{k+1}(K_{k+1})$, which is not needed (and very expensive)
- The resulting persistence barcode is the same
[de Silva et al. 2011]

Persistent cohomology

We have seen: many columns of $R = D \cdot V$ are not needed

- Skip those inessential columns in matrix reduction

For persistence barcodes in low dimensions $d \leq k$:

- Number of skipped indices for reducing D^T (cohomology) is much larger than for D (homology)
 - reducing boundary matrix produces basis for $H_{k+1}(K_{k+1})$, which is not needed (and very expensive)
- The resulting persistence barcode is the same
[de Silva et al. 2011]

Example ($k = 2, n = 192$):

- $\underbrace{1161\,471}_{\text{negative}} + \underbrace{1161\,471}_{\text{skipped}} + \underbrace{53\,727\,345}_{\text{essential}}$ columns in homology
- $\underbrace{1161\,471}_{\text{negative}} + \underbrace{18\,336}_{\text{skipped}} + \underbrace{1}_{\text{essential}}$ columns in cohomology

Observations

For a typical input:

- V has very few off-diagonal entries
- most negative columns of D are already reduced

Observations

For a typical input:

- V has very few off-diagonal entries
- most negative columns of D are already reduced

Previous example ($k = 2, n = 192$):

- Only 845 out of 1 161 471 columns have to be reduced

Implicit matrix reduction

Implicit matrix reduction

Standard approach:

- Boundary matrix D for filtration-ordered basis
 - Explicitly generated and stored in memory

Implicit matrix reduction

Standard approach:

- Boundary matrix D for filtration-ordered basis
 - Explicitly generated and stored in memory
- Matrix reduction: store only reduced matrix R
 - transform D into R by column operations

Implicit matrix reduction

Standard approach:

- Boundary matrix D for filtration-ordered basis
 - Explicitly generated and stored in memory
- Matrix reduction: store only reduced matrix R
 - transform D into R by column operations

Approach for Ripser:

- Boundary matrix D for lexicographically ordered basis
 - Implicitly defined and recomputed when needed

Implicit matrix reduction

Standard approach:

- Boundary matrix D for filtration-ordered basis
 - Explicitly generated and stored in memory
- Matrix reduction: store only reduced matrix R
 - transform D into R by column operations

Approach for Ripser:

- Boundary matrix D for lexicographically ordered basis
 - Implicitly defined and recomputed when needed
- Matrix reduction in Ripser: store only coefficient matrix V
 - recompute previous columns of $R = D \cdot V$ when needed
 - Typically, V is much sparser and smaller than R

Oblivious matrix reduction

Algorithm variant:

- $R = D$
- for $j = 1, \dots, n$
 - while $\exists i < j$ with $\text{pivot } R_i = \text{pivot } R_j$
 - add D_i to R_j

Oblivious matrix reduction

Algorithm variant:

- $R = D$
- for $j = 1, \dots, n$
 - while $\exists i < j$ with $\text{pivot } R_i = \text{pivot } R_j$
 - add D_i to R_j

Computing the persistence intervals requires only:

- current column R_j
- pivots of previous columns R_i

Oblivious matrix reduction

Algorithm variant:

- $R = D$
- for $j = 1, \dots, n$
 - while $\exists i < j$ with $\text{pivot } R_i = \text{pivot } R_j$
 - add D_i to R_j

Computing the persistence intervals requires only:

- current column R_j
- pivots of previous columns R_i

Corollary

The rank of an $m \times n$ matrix can be computed in $O(n)$ memory.

Apparent and emergent pairs

Natural filtration settings

Typical assumptions on the filtration:

- | | |
|-------------------------------------|---------------------------|
| • general filtration | persistence (in theory) |
| • filtration by singletons or pairs | discrete Morse theory |
| • simplexwise filtration | persistence (computation) |

Natural filtration settings

Typical assumptions on the filtration:

- | | |
|-------------------------------------|---------------------------|
| • general filtration | persistence (in theory) |
| • filtration by singletons or pairs | discrete Morse theory |
| • simplexwise filtration | persistence (computation) |

Conclusion:

- Discrete Morse theory sits in the middle between persistence and persistence (!)

Apparent pairs

Definition

In a simplexwise filtration, (σ, τ) is an *apparent* pair if

- σ is the youngest face of τ
- τ is the oldest coface of σ

Apparent pairs

Definition

In a simplexwise filtration, (σ, τ) is an *apparent* pair if

- σ is the youngest face of τ
- τ is the oldest coface of σ

Lemma

Any apparent pairs is a persistence pair.

Lemma

The apparent pairs form a discrete gradient.

- Generalizes a construction proposed by [Kahle 2011] for the study of random Rips filtrations

Emergent persistent pairs

Consider the *lexicographically refined Rips filtration*:

- increasing diameter, refined by
- lexicographic order

This is the simplexwise filtration for computations in Ripser.

Emergent persistent pairs

Consider the *lexicographically refined Rips filtration*:

- increasing diameter, refined by
- lexicographic order

This is the simplexwise filtration for computations in Ripser.

Lemma

Assume that

- τ is the lexicographically minimal proper coface of σ with $\text{diam}(\tau) = \text{diam}(\sigma)$,
- and τ is not already in a persistence pair (ρ, τ) with $\rho > \sigma$.

Then (σ, τ) is an emergent persistence pair.

Emergent persistent pairs

Consider the *lexicographically refined Rips filtration*:

- increasing diameter, refined by
- lexicographic order

This is the simplexwise filtration for computations in Ripser.

Lemma

Assume that

- τ is the lexicographically minimal proper coface of σ with $\text{diam}(\tau) = \text{diam}(\sigma)$,
- and τ is not already in a persistence pair (ρ, τ) with $\rho > \sigma$.

Then (σ, τ) is an emergent persistence pair.

- Includes all apparent pairs with persistence 0
- Can be identified *without* enumerating all cofaces of σ
 - Provides a shortcut for computation

Ripsers Live: users from 156 different cities

